

ALEX MORGAN

Email: alex.morgan.ai@example.com | Phone: (555) 321-7654 | San Francisco, CA

LinkedIn: linkedin.com/in/alexmorgan-ai | GitHub: github.com/alexmorgan-ai

PROFESSIONAL SUMMARY

Senior Machine Learning Engineer with 4+ years of experience specializing in NLP, Deep Learning, and Transformer architectures. Proven track record deploying semantic search engines, fine-tuning LLMs, and architecting low-latency microservices with PyTorch, FastAPI, and Docker on AWS.

TECHNICAL SKILLS

- Programming Languages: Python, SQL, Bash, C++
- AI & Data Science: Machine Learning, Deep Learning, Natural Language Processing (NLP), Large Language Models (LLMs), Transformers, PyTorch, Scikit-Learn, Hugging Face, Pandas, NumPy
- Cloud & DevOps: Docker, AWS (EC2, S3, ECS), Git, CI/CD, Linux
- Databases & Backend: FastAPI, RESTful APIs, PostgreSQL, Redis, SQLite
- Practices: Agile, Scrum, Unit Testing, System Design

PROFESSIONAL EXPERIENCE

Senior AI Engineer | Apex Intelligence Systems | 2022 - Present

- Architected enterprise semantic document retrieval pipeline using Sentence-Transformers and vector embeddings, improving search precision by 42%.
- Fine-tuned transformer models for sentiment analysis and entity extraction handling 3M+ daily inference queries.
- Containerized NLP inference microservices with Docker and deployed on AWS ECS with auto-scaling, reducing p99 latency to under 45ms.
- Built automated data cleaning and feature engineering pipelines using Pandas, NumPy, and Scikit-Learn.

Machine Learning Engineer | DataVision Labs | 2020 - 2022

- Developed text classification and intent recognition models with PyTorch and Scikit-Learn achieving 94.6% F1-score.
- Created RESTful APIs using FastAPI and integrated Redis caching for prompt management.
- Collaborated in cross-functional Agile sprints with product managers and frontend developers.

EDUCATION

B.S. in Computer Science | University of California, Berkeley | 2016 - 2020

- Focus: Artificial Intelligence & Computational Linguistics